

Neuromorphic Intermediate Representation (NIR)

A common language for neuromorphic systems

Jens E. Pedersen · jegpe@dtu.dk


DTU Electro

ISCAS 2026 Tutorial — SNNs to Silicon

The problem: a fragmented ecosystem

Training an SNN today means choosing a stack:

- Framework: Nornox, snnTorch, Lava-DL, Sinabs, Spyx, ...
- Hardware: Loihi, BrainScaleS, SpiNNaker, Speck, Xylo, FPGA, ...
- Each combination requires **rewriting** the model




Pedersen et al., Nature Reviews Electrical Engineering, 2025

The problem: a fragmented ecosystem

Training an SNN today means choosing a stack:

- Framework: Nornox, snnTorch, Lava-DL, Sinabs, Spyx, ...
- Hardware: Loihi, BrainScaleS, SpiNNaker, Speck, Xylo, FPGA, ...
- Each combination requires **rewriting** the model



This is the same problem conventional ML solved with **ONNX** — but neuromorphic models have **dynamics** that ONNX cannot capture.



Pedersen et al., Nature Reviews Electrical Engineering, 2025

NIR: one model, many platforms



NIR captures SNN models as graphs of continuous-time dynamical systems



NIR: one model, many platforms

NIR captures SNN models as graphs of continuous-time dynamical systems



- Primitives are parameterized ODEs — no discretization baked in



NIR: one model, many platforms

NIR captures SNN models as graphs of continuous-time dynamical systems



- Primitives are parameterized ODEs — no discretization baked in
- The graph is substrate-agnostic



NIR: one model, many platforms

NIR captures SNN models as graphs of continuous-time dynamical systems

- Primitives are **parameterized ODEs** — no discretization baked in
- The graph is **substrate-agnostic**
- Each backend instantiates the dynamics:
 - Software → any ODE solver or Euler step
 - Digital HW → fixed-point discretization
 - Analog HW → physics implements the ODE
 - FPGA → streaming dataflow (NIR2FPGA)



What's in a NIR graph?

17 primitives covering the neuromorphic toolbox:

Category	Primitives
Neurons	LIF, LI, IF, CubaLIF, CubaLI
Synapses	Linear, Affine, Conv1d, Conv2d
Pooling	SumPool2d, AvgPool2d
Utility	Scale, Delay, Flatten, Threshold
I/O	Input, Output

What's in a NIR graph?

17 primitives covering the neuromorphic toolbox:

Category	Primitives
Neurons	LIF, LI, IF, CubaLIF, CubaLI
Synapses	Linear, Affine, Conv1d, Conv2d
Pooling	SumPool2d, AvgPool2d
Utility	Scale, Delay, Flatten, Threshold
I/O	Input, Output

The LIF neuron in NIR:

$$\tau \frac{dv}{dt} = -(v - v_{\text{leak}}) + R \cdot I$$

if $v \geq \theta$: spike, $v \leftarrow v_{\text{reset}}$

Parameters: τ , R , v_{leak} , v_{reset} , θ

→ Fully defines the dynamics, independent of timestep or solver

Supported platforms

10 software frameworks

- Norse (PyTorch)
- snnTorch (PyTorch)
- Sinabs (PyTorch / SynSense)
- Lava-DL (Intel)
- Rockpool (SynSense)
- Nengo
- Spyx (JAX)
- jaxsnn (BrainScaleS)
- hxtorch (BrainScaleS)
- Spiking Jelly

5+ hardware targets

- Intel Loihi
- SynSense Speck & Xylo
- BrainScaleS-2
- SpiNNaker2
- FPGA via NIR2FPGA ← this tutorial!


Supported platforms

10 software frameworks

- Norse (PyTorch)
- snnTorch (PyTorch)
- Sinabs (PyTorch / SynSense)
- Lava-DL (Intel)
- Rockpool (SynSense)
- Nengo
- Spyx (JAX)
- jaxsnn (BrainScaleS)
- hxtorch (BrainScaleS)
- Spiking Jelly

5+ hardware targets

- Intel Loihi
- SynSense Speck & Xylo
- BrainScaleS-2
- SpiNNaker2
- FPGA via NIR2FPGA ← this tutorial!



NIR in practice: export and import

Export from any supported framework:

```
import nir
import norse

# Train your model in Norse
net = norse.torch.SequentialState( ... )
# ... training loop ...

# Export to NIR
nir_graph = norse.torch.to_nir(net)
nir.write("my_snn.nir", nir_graph)
```

Import into any target:

```
import nir

# Load the NIR graph
graph = nir.read("my_snn.nir")

# Inspect the model
for name, node in graph.nodes.items():
    print(f"{name}: {type(node).__name__}")
    # e.g., "lif_0: LIF"
    # e.g., "linear_0: Affine"

# Deploy to hardware, FPGA, or simulator
```

NIR in practice: export and import

Export from any supported framework:

```
import nir
import norse

# Train your model in Norse
net = norse.torch.SequentialState( ... )
# ... training loop ...

# Export to NIR
nir_graph = norse.torch.to_nir(net)
nir.write("my_snn.nir", nir_graph)
```

Import into any target:

```
import nir

# Load the NIR graph
graph = nir.read("my_snn.nir")

# Inspect the model
for name, node in graph.nodes.items():
    print(f"{name}: {type(node).__name__}")
    # e.g., "lif_0: LIF"
    # e.g., "linear_0: Affine"

# Deploy to hardware, FPGA, or simulator
```

File format is **HDF5** — portable, self-describing, and language-agnostic

Building a NIR graph from scratch

```
import nir
import numpy as np

graph = nir.NIRGraph(
    nodes={
        "input": nir.Input({"input": np.array([784])}),
        "linear": nir.Affine(weight=np.random.randn(128, 784),
                             bias=np.zeros(128)),
        "lif": nir.LIF(tau=np.full(128, 20.0),
                      r=np.ones(128),
                      v_leak=np.zeros(128),
                      v_threshold=np.ones(128),
                      v_reset=np.zeros(128)),
        "out": nir.Affine(weight=np.random.randn(10, 128),
                          bias=np.zeros(10)),
        "output": nir.Output({"output": np.array([10])}),
    },
    edges=[("input", "linear"), ("linear", "lif"),
           ("lif", "out"), ("out", "output")]
)
nir.write("mnist_snn.nir", graph) # Ready for NIR2FPGA!
```

synfire.dev: the neuromorphic model registry

An open platform for sharing and deploying SNN models

synfire.dev: the neuromorphic model registry

An open platform for sharing and deploying SNN models

- Model registry: publish and discover SNNs

synfire.dev: the neuromorphic model registry

An open platform for sharing and deploying SNN models

- Model registry: publish and discover SNNs
- Built on NIR: models are stored as graphs

synfire.dev: the neuromorphic model registry

An open platform for sharing and deploying SNN models

- Model registry: publish and discover SNNs
- Built on NIR: models are stored as graphs
- Hardware-aware deployment: deploy to neuromorphic chips and FPGAs

synfire.dev: the neuromorphic model registry

An open platform for sharing and deploying SNN models

- Model registry: publish and discover SNNs
- Built on NIR: models are stored as graphs
- Hardware-aware deployment: deploy to neuromorphic chips and FPGAs
- CLI + SDK + Web: multiple developer interfaces

synfire.dev: the neuromorphic model registry

An open platform for sharing and deploying SNN models

- **Model registry:** publish and discover SNNs
- **Built on NIR:** models are stored as graphs
- **Hardware-aware deployment:** deploy to neuromorphic chips and FPGAs
- **CLI + SDK + Web:** multiple developer interfaces

Think of it as HuggingFace for neuromorphic AI — train once, deploy everywhere

synfire.dev: the neuromorphic model registry

An open platform for sharing and deploying SNN models

- **Model registry:** publish and discover SNNs
- **Built on NIR:** models are stored as graphs
- **Hardware-aware deployment:** deploy to neuromorphic chips and FPGAs
- **CLI + SDK + Web:** multiple developer interfaces


Think of it as **HuggingFace** for neuromorphic AI — train once, deploy everywhere

Workflow:

1. Train SNN in your preferred framework
2. Export to NIR
3. Upload to synfire.dev
4. Others discover & deploy your model
5. Target any supported hardware

The NIR2FPGA pipeline


This tutorial's end-to-end flow:



The NIR2FPGA pipeline

This tutorial's end-to-end flow:


1. **Define an SNN (JAX — next session)**



The NIR2FPGA pipeline

This tutorial's end-to-end flow:


1. **Define** an SNN (JAX — next session)
2. **Train** with surrogate gradients



The NIR2FPGA pipeline

This tutorial's end-to-end flow:


1. **Define** an SNN (JAX — next session)
2. **Train** with surrogate gradients
3. **Export** to NIR



The NIR2FPGA pipeline

This tutorial's end-to-end flow:


1. **Define** an SNN (JAX — next session)
2. **Train** with surrogate gradients
3. **Export** to NIR
4. **Compile** NIR → SpinalHDL RTL



The NIR2FPGA pipeline

This tutorial's end-to-end flow:

1. **Define** an SNN (JAX — next session)
2. **Train** with surrogate gradients
3. **Export** to NIR
4. **Compile** NIR → SpinalHDL RTL
5. **Simulate** or **synthesize** for FPGA




The NIR2FPGA pipeline

This tutorial's end-to-end flow:

1. **Define** an SNN (JAX — next session)
2. **Train** with surrogate gradients
3. **Export** to NIR
4. **Compile** NIR → SpinalHDL RTL
5. **Simulate** or **synthesize** for FPGA

NIR closes the gap between **software simulation** and **hardware realization** — same model, verified at every step



Summary & resources

NIR is the common language for neuromorphic systems:

- 17 ODE-based primitives
- 10 frameworks, 5 (now 6!) hardware targets
- HDF5 file format, open source
- synfire.dev for model sharing
- NIR2FPGA for FPGA deployment

References & links

- neuroir.org — NIR documentation
- synfire.dev — model registry
- github.com/neuromorphs/NIR
- Paper: [Neuromorphic Intermediate Representation](#), Nat. Commun, 2025
- Paper: [Neuromorphic Programming](#), ICONS, 2024
- snnbook.net — free SNN textbook

Next up: Defining, training & exporting SNNs →